

Master Theorem

$$T(n) = aT\left(\frac{n}{b}\right) + f(n), \quad f(n) = \Theta(n^k \log^p n), \text{ where } a \geq 1 \text{ and } b > 1$$

Case 1: if $\log_b a > k$

- then, $\Theta(n^{\log_b a})$

Case 2: if $\log_b a = k$

- if $p > -1$, then $\Theta(n^k \log^{p+1} n)$
- if $p = -1$, then $\Theta(n^k \log \log n)$
- if $p < -1$, then $\Theta(n^k)$

Case 3: if $\log_b a < k$

- if $p \geq 0$, then $\Theta(n^k \log^p n)$
- if $p < 0$, then $\Theta(n^k)$

Master Theorem for subtract and conquer

$$T(n) = aT(n-b) + f(n), \quad n > 1 \text{ for } a > 0, b > 0, f(n) = O(n^k), k \geq 0$$

Case 1: if $a < 1$, then
 $T(n) = O(n^k)$

Case 2: if $a = 1$, then
 $T(n) = O(n^{k+1})$

Case 3: if $a > 1$, then
 $T(n) = O(n^k a^{\frac{n}{b}})$

Recurrence relation	big-Oh	Recurrence relation	big-Oh
$T(n) = T(n/2) + O(1)$	$O(\log n)$	$T(n) = 2T(n-1) - 1$	$O(1)$
$T(n) = T(n/2) + O(n)$	$O(n)$	$T(n) = T(\sqrt{n}) + O(1)$	$O(\log \log n)$
$T(n) = T(n/2) + O(n^2)$	$O(n^2)$	$T(n) = 2T(\sqrt{n}) + O(1)$	$O(\log n)$
$T(n) = T(n-1) + O(1)$	$O(n)$	$T(n) = T(\sqrt{n}) + O(n)$	$\approx O(n)$
$T(n) = 2T(n/2) + O(1)$	$O(n)$	$T(n) = 2T(\sqrt{n}) + O(n)$	$O(n)$
$T(n) = T(n-1) + O(n)$	$O(n^2)$	$T(n) = 2T(n-1) + 1$	$O(2^n)$
$T(n) = 2T(n/2) + O(n)$	$O(n \log n)$	$T(n) = 3T(n-1)$	$O(3^n)$

Sorting Algorithms - Best | Average | Worst | Stability

Bubble: $\Omega(N)|\Theta(N^2)|O(N^2)|\text{Stable}$ → At the end of iteration j , the biggest j items are correctly sorted in the right most j position of array.

Selection: $\Omega(N^2)|\Theta(N^2)|O(N^2)|\text{Not Stable}$ → At the end of iteration j : the **smallest** j items are correctly sorted in the first j positions of the array. Always traverse entire array, always N^2 .

Insertion: $\Omega(N)|\Theta(N^2)|O(N^2)|\text{Stable}$ → At the end of j iteration: the first j items in the array are in sorted order. (**NOT** necessarily the smallest). Best/Worst: Already sorted vs Reverse Sorted

Quick: $\Omega(N \log N)|\Theta(N \log N)|O(N^2)|\text{Stable}$ → At the end of every loop iteration: for all $i \geq \text{high}$, $A[i] > \text{pivot}$, for all $1 < j < \text{low}$, $A[j] < \text{pivot}$. Worst: Pivot smallest/biggest item.

Merge: $\Omega(N \log N)|\Theta(N \log N)|O(N \log N)|\text{Stable}$ → At the end of j iteration, each half of the array that is split is already sorted. Stable if `merge()` is stable.

Heap: $\Omega(N \log N)|\Theta(N \log N)|O(N \log N)$

Binary Search: $\Omega(1)|\Theta(\log N)|O(\log N)$ → At the end of iteration j , the active search region between begin and end is size at most $\frac{n}{2^j}$, and the item (if it exists) is in that region.

```
static int genBinarySearch(int[] arr, int key, int k) {
    int low = 0; int high = arr.length - 1;
    while(low <= high) {
        int pivot = low + (high - low) / k; //two halves split with fraction k
        if (arr[pivot]==key) return pivot; //corner case: found target
        else if (arr[pivot] > key) { high = pivot - 1; } //decrease high (search LHS)
        else { low = pivot + 1; } //increase low (search RHS)
    } return -1 //not found anywhere;
} //realise that when k=2, this reduces to our vanilla binary search
```

Data Structures

1. Binary Trees - Unbalanced, in-order, dynamic, traversable

Time complexity of all BST Operations = $O(h)$, where h is the height of BST.

Invariant: all in left sub-tree < key < all in right sub-tree.

Best Case: BST is balanced, height of BST is $O(\log n)$, hence complexity is $O(\log n)$.

Worst Case: BST is skewed, height of BST becomes n , time complexity is $O(n)$.

2. Tries - Unbalanced, string-based

Best Case/Worst Case: For all operations of Insertion/Deletion/Searching $O(N)$ except if the word is 1 letter long, in which case search time is $O(1)$ (best case for searching).

3. (a,b) trees - Degree [a,b], balanced, uniform height

Invariants: After every operation, every leaf has same depth. Every node has degree in the range $[a,b]$. Nodes satisfy tree-order property.

Node	Min Keys	Max Keys	Min Child	Max Child
Root	1	$b-1$	2	b
Internal	$a-1$	$b-1$	a	b
Leaf	$a-1$	$b-1$	0	0

- An (a, b) -tree with n nodes has $O(\log_a n)$ height, binary search for a key at every node takes $O(\log_2 b)$ time.

- Total cost = $O(\log_2 b * \log_a n)$
(Cost of searching a node * Height)
= $O(\log_a n) * (\log_a b \text{ const}) = O(\log_a n)$

4. AVL Trees - Like binary trees, but height-balanced. Height-balanced tree has at most height $h < 2\log(n)$, \therefore at least $n > 2^{\frac{h}{2}}$ nodes. at most $2^{h+1} - 1$ nodes.

Invariants: After every operation, every node is height-balanced.

Best Case: $O(\log n)$, for insertion/deletion/searching.

Insertions take at most 2 rotations to balance as insertions DO NOT reduce heights. Deletion can take more than 2 rotations to balance. Rebalancing does not “undo” the change in height since rebalances usually reduces height as well.

Worst Case: should tree be unbalanced from leaf-to-root upon deletions, it would take $\log n$ rotations still, hence worst time complexity would be $O(\log n)$.

Updating AVL Trees:

- Rotations:** Rotations only affect current rotating node and parents of rotating node, thus we update from root-to-currNode path.
- Inserting:** Depend on how the variable affected is initialized, update while inserting/after.
 - If balanced, update variables from root-to-leaf path (if needed).
 - If not balanced, update first and conduct necessary rotations.
- Deleting:** Three cases
 - Node is just a leaf → just delete and update parents.
 - 1 child → just delete and update from root to deleted node's variable, connect **parent** of deleted node to **child** of deleted node.
 - 2 children → Find successor (maximum of left-subtree of deleted node), delete node, replace with successor and update variables. Have to update the parents of deleted nodes and the left sub-tree of deleted node (since a node is taken from there)

5. Interval Trees Search(key) $\Rightarrow O(\log n)$

If value is in root, return. If value > max(left sub-tree) recurse right. else left (can't go right).

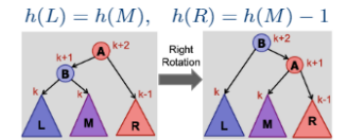
6. Orthogonal Range Searching: Binary trees; leaves store points internal nodes ← left subtree max

$$\text{G.P } S_n = \frac{a(1-r^n)}{1-r}$$

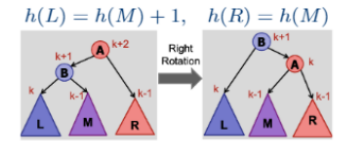
$$\text{A.P } S_n = \frac{n}{2}[2a + (n-1)d] \text{ OR } S_n = \frac{n}{2}[a + l]$$

rebalancing

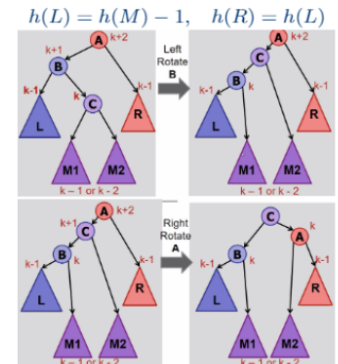
[case 1] B is **balanced**: **right-rotate**



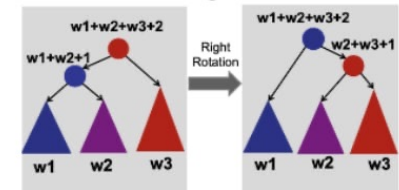
[case 2] B is **left-heavy**: **right-rotate**



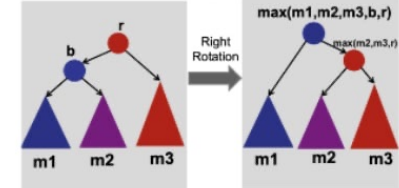
[case 3] B is **right-heavy**: **left-rotate(v.left)**, **right-rotate(v)**



weights



max



Series	Asymptotic Bound
$n + \frac{n}{2} + \frac{n}{3} + \dots + \frac{n}{n}$ $= n(1 + \frac{1}{2} + \dots + \frac{1}{n})$	$\approx n(\log n)$ $O(n \log n)$
$1 + 2 + 3 + 4 + \dots + n = \frac{n(n+1)}{2}$	$= O(n^2)$
$n + \frac{n}{2} + \dots + 1$ $= n + n(\frac{1}{2}) + n(\frac{1}{2})^2 + \dots + n(\frac{1}{2})^n$	$\approx O(n)$
$1 + c + c^2 + \dots + c^n$ if $ c < 1 = O(1)$ OR if $ c > 1 = O(c^n)$	
$1^2 + 2^2 + 3^2 + \dots + n^2 = \frac{n(n+1)(2n+1)}{6}$	$O(n^3)$

7. Hash Tables

Let m be the table size and let n be the number of items; let $cost(h)$ be cost of hash function. $load(hash\ table), \alpha = \frac{n}{m} \rightarrow$ average/expected number of items per bucket.

- **Simple Uniform Hashing Assumption:** every key has an equal probability of being mapped to every bucket, keys are mapped independently.
- **Uniform Hashing Assumption:** every key is equally likely to be mapped to every permutation, independent of every other key, NOT fulfilled by linear probing.
- **Properties of a good hash function:** able to enumerate all possible buckets - $h : U \rightarrow \{1...m\}$ for every bucket $j, \exists i$ such that $h(key, i) = j$ && SUHA

Rules of `hashCode()` function: (rules for equals method \rightarrow equivalence relation).

1. Always returns the same value if the object has not changed.
2. If two objects are equal, then they return the same hashCode.
3. MUST redefine `.equals()` to be consistent with `hashCode()`.

Chaining

`insert(key, value) $\rightarrow O(1 + cost(h)) \Rightarrow O(1)$,`
for n items, expected max cost = $O(logn) \rightarrow O(\frac{logn}{log(logn)})$

`search(key) \rightarrow Worst case: $O(n + cost(h)) \Rightarrow O(n)$,`
expected case, $O(\frac{n}{m} + cost(h)) \Rightarrow O(1)$ (Searching through a linked list).

Open addressing - linear probing

redefined hash function: $h(k, i) = h(k, 1) + i \mod m$

`delete(key)` : use a tombstone value - DON'T set to `null`

Performance

if the table is $\frac{1}{4}$ full, there will be clusters of size $\Theta(logn)$ expected cost of an operation, $E[\#probes] \leq \frac{1}{1-\alpha}$, (assume $\alpha < 1$ and uniform hashing), degrades badly as $\alpha \rightarrow 1$

advantages

- saves space (use empty slots vs linked list)
- better cache performance (table is one place in memory)
- rarely allocate memory (no new list-node allocation)

disadvantages

- more sensitive to choice of hash function (clustering)
- more sensitive to load (as $\alpha \rightarrow 1$, performance degrades)

8. Table size

- If $(n == m)$, then $m = 2m$.
If $(n < \frac{m}{4})$ then $m = \frac{m}{2}$.
- Everytime double a table of size m , at least $\frac{m}{2}$ new items were added (pay for doubling).
- Everything shrink a table of size m , at least $\frac{m}{4}$ items were deleted (pay for shrinking).

- Let m_1 be size of old hash table, m_2 be size of new hash table; n be number of elements. growing the table:
 $O(m_1 + m_2 + n) \approx O(n)$

table growth	resize	insert n items
increment by 1	$O(n)$	$O(n^2)$
double	$O(n)$	$O(n)$, average $O(1)$
square	$O(n^2)$	$O(n)$

9. Amortized Analysis - an operation has amortized cost $T(n)$ if for every integer k , the cost of k operations is $\leq kT(n)$.

- binary counter ADT: increment $\Rightarrow O(1)$
- hash table resizing: $O(k)$ for k insertions, search operation: expected $O(1)$ (not amortized)

$$1 < \log n < \sqrt{n} < n < n \log n < n^2 < n^3 < 2^n < 2^{2n}$$
$$\log_a n < n^a < a^n < n! < n^n$$

10. Graphs

- diameter: max. shortest path in the graph.
 - even cycles are bipartite (every other node 1 grp).
 - degree (graph): maximum degree of a node.
 - Searching
- degree (node): number of adjacent edges
 - graph is **dense** if $|E| = \Theta(V^2)$, if **dense** use a Adjacency matrix instead.

adj	space	(cycle)	(clique)	use for
list	$O(V + E)$	$O(V)$	$O(V^2)$	sparse
matrix	$O(V^2)$	$O(V^2)$	$O(V^2)$	dense

BFS $\Rightarrow O(V + E)$, **parent edges** form a **tree** & shortest path from S , use queue.

- $O(V)$ - every vertex is added exactly once to a frontier
- $O(E)$ - every neighbourList is enumerated once

DFS $\Rightarrow O(V + E)$, with adjacency matrix: $O(V)$ per node \Rightarrow total $O(V^2)$, use stack

- $O(V)$ - DFSvisit is called exactly once per node
- $O(E)$ - DFSvisit enumerates each neighbour

• Topological Ordering - Sequential total ordering of nodes + Edges only point forward.

- Post-Order DFS $\Rightarrow O(V + E)$, prepend each node from a post-order traversal.
- Khan's Algorithm $\Rightarrow O(E \log V)$, add nodes without incoming edges to topological order.
 - remove min-degree node from priority queue $\Rightarrow O(V \log V)$
 - `decreaseKey()` (in degree) of the child node $\Rightarrow O(E \log V)$

• Shortest Paths

- **Bellman Ford** $\Rightarrow O(VE)$, $|V|$ iterations of relaxing, terminate when entire sequence of $|E|$ operations have no effect. Use to detect negative weight cycles, if same weight use BFS.
- **Dijkstra's** $\Rightarrow O((E + V) \log V) = O(E \log V) \rightarrow$ NO Negative weight edges.

Dijkstra grows set by adding neighbouring node that is the **closest to source** (min-estimate)
Use a PQ to track the min-estimate to node, relax outgoing edges, add incoming to PQ.

$|V|$ times of `insert()` / `deleteMin()` $\rightarrow O(\log V)$ each, (only inserted/removed once)

$|E|$ times of `relax()` / `decreaseKey()` $\rightarrow O(\log E)$ each.

Fibonacci heap $\Rightarrow O(E + V \log V)$, Array $\Rightarrow O(V^2)$

- **DAG Shortest Path** $\Rightarrow O(V + E)$, Sort by topological order, then relax in this order
Longest path: negate the edges/modify relax function
- **Trees** $\Rightarrow O(V)$, relax each edge in BFS/DFS order.

• Union-Find, Union - connect 2 objects, Find - check if objects are connected.

	Find	Union
quick-find <code>int[] componentId</code> flat trees (connected if part of same componentId)	$O(1)$ - check if objects have the same <code>componentId</code>	$O(n)$ - enumerate all items in array to update <code>id</code>
quick-union <code>int[] parent</code> deeper trees (connected if same parent)	$O(n)$ - check for same root	$O(n)$ - add as a subtree of the root
weighted union <code>int[] parent, int[] size</code> (make shorter tree child of larger tree)	$O(\log n)$ - check for same root	$O(\log n)$ - add as smaller subtree of root
path compression (set parent of each traversed node to the root)	$O(\log n)$	$O(\log n)$
weighted union + path compression , for m union/find operations on n objects: $O(n + m\alpha(m, n))$, flat trees	$O(\alpha(m, n))$	$O(\alpha(m, n))$

11. Minimum Spanning Trees

Property 1: No cycles

Property 2: If you cut an MST, the two subtrees are both MSTs.

Property 3: Cycle property, For every cycle, the max weight edge is **not** in the MST.

False cycle property, For every cycle, the min weight edge may not be in the MST.

Property 4: Cut property, For every partition of nodes, min weight edge across cut is in the MST.
 \rightarrow for every vertex, minimum outgoing edge is in the MST.

Prim's Algorithm $\Rightarrow O(E \log V)$, grows set by adding node connected via lightest edge.

- add the minimum edge across the cut in the MST
- PQ to store nodes (priority: lowest incoming edge weight)
- each vertex: one `insert()` / `extractMin()` $\Rightarrow O(\log V)$ each, hence total $O(V * \log V)$
- each edge: one `decreaseKey()` $\Rightarrow O(E \log V)$

Kruskal's Algorithm $\Rightarrow O(E \log V)$, sort edges, add smallest edges that do not form a cycle.

• sorting $\Rightarrow O(E \log E) = O(E \log V)$

• each edge: `find()` / `union()` $\Rightarrow O(\log V)$ using union-find DFS

12. Dynamic Programming

Optimal sub-structure - optimal solution constructed from solutions to smaller sub-problems

Overlapping sub-problems - can memoize solutions to sub problems

- optimal substructure but no overlapping sub-problems \rightarrow divide-and-conquer

• Longest Increasing sub-sequence

- DAG Solution $\Rightarrow O(n^3)$, find topological order ($O(V + E) = O(n^2)$), ran n times
- $S[i] = \text{LIS}(A[1...i])$, $\Rightarrow O(n^2)$, find LIS up to each point in $A[1...i] = O(n)$, ran n times
- **Prize collecting** $\Rightarrow O(kE) \rightarrow$ super source connects k graphs $|O(kV^2)|$ for k steps (DP)
- **Vertex cover** $\Rightarrow O(V)$ or $O(V^2)$, set of nodes where every edge is adjacent to ≥ 1 node
- **All pairs shortest path** $\Rightarrow O(V * E \log V)$, dijkstra ran onto all nodes once (V times)
- **Diameter of a graph:** $O(V^2 \log V)$, SSSP all nodes of a graph, (V times)
- **Floyd Warshall** $\Rightarrow O(V^3)$

- $S'[v, w, P_k]$ = shortest path from v to w only using nodes from set P
- $S'[v, w, P_8] = \min(S'[v, w, P_7], S'[v, 8, P_7] + S'[8, w, P_7])$

