

Database Management System

Transactions

defined as a **finite sequence** of database operations *read/write* and constitutes the smallest logical unit of work from the application perspective.

Properties of each transaction, T

- Atomiticty - all/none of T reflected in the database
- Consistency - T guarantees the correct state
- Isolation - T isolated from effects of concurrent transactions
- Durability - after T, effects are permanent

Definitions

- A *data model* is a collection of concepts for describing data.
- A *schema* is a description of the structure of a database using a data model.

Relational Data Model

A relational database schema is a set of relation schemas + data constraints. Data is modelled by *relations*, each *relation* defined as a **relation schema**:

- attributes: columns
- data constraints: domain constraints → datatype
- `Employees(id: INT, name: TEXT, salary: NUMERIC())`

A relation can be seen as a *table with rows and columns*

- Number of columns = *Degree/Arity*
- Number of rows = *Cardinality*

A **domain** is a set of atomic values (INT, NUMERIC, TEXT) denoted as *dom(A<sub>i</sub>)*. A **relation** is a set of tuples (records in the database), denoted as *R(A<sub>1</sub>, A<sub>2</sub>, ..., A<sub>n</sub>)*.  
• every v of attribute A<sub>i</sub> is either: v ∈ *dom(A<sub>i</sub>)* / v = NULL.  
• we write *R.A<sub>i</sub>* to refer to the attribute A<sub>i</sub> of relation R.

Data Integrity

defined as conditions that restrict what constitutes valid data.

Keys

- **superkey**: subset of attributes that uniquely identifies a tuple in a relation
- **key**: is a *superkey* that is also minimal → no proper subset of the key is a superkey.
  - minimal (cannot be made smaller) ≠ minimum (smallest)
- **candidate keys**: set of all keys of a given relation
- **primary key**: selected candidate keys, and they CANNOT be NULL. To simplify our notation, we underline our primary keys in the schema notation.

Foreign keys

- subset of attributes of relation R<sub>1</sub> that refers to the PK of relation R<sub>2</sub>. → R<sub>1</sub>: referencing, R<sub>2</sub>: referenced relation
- **requirements**: each FK in R<sub>1</sub> must either appear as a PK in R<sub>2</sub> **OR** be a NULL value (in a tuple, containing at least 1 NULL value)

Relational Algebra

Closure Property

A set of values is **closed** under the set of operators if any combinations of the operators produces values in the given set  
• Relations *closed* under RA, we can chain operations

Three-Valued Logic

Logical Operations

Truth Table

c <sub>1</sub>	c <sub>2</sub>	c <sub>1</sub> ∧ c <sub>2</sub>	c <sub>1</sub> ∨ c <sub>2</sub>	¬c <sub>1</sub>
False	False	False	False	True
False	NULL	False	NULL	True
False	True	False	True	True
NULL	False	False	NULL	NULL
NULL	NULL	NULL	NULL	NULL
NULL	True	NULL	True	NULL
True	False	False	True	False
True	NULL	NULL	True	False
True	True	True	True	False

Basic Operators

Selection: σ<sub>c</sub>

σ<sub>c</sub>(R) selects tuples from relation R, satisfying the selection condition c. e.g.: σ<sub>price<20</sub>(Sells). The *selection condition* is a boolean condition of *terms*. A *term* is one of the following forms:

*attribute op constant; attribute<sub>1</sub> op attribute<sub>2</sub>; term<sub>1</sub> and term<sub>2</sub>; term<sub>1</sub> or term<sub>2</sub>; not term<sub>1</sub>; (term<sub>1</sub>)*

- op ∈ {=, <>, <, ≤, >, ≥}
- Operator precedence: (), op, not, and, or.
- A tuple is only selected if condition evaluates to *true* on it.

Projection ρ<sub>l</sub>

l is a list of attribute renamings of the form a<sub>1</sub>:b<sub>1</sub>, ..., a<sub>n</sub>:b<sub>n</sub>. The order of the attribute renamings in l does not matter, however it will remove any duplicates after projection.

Set Operators

The attributes require input relations to be *union compatible*:

*Union*: R ∪ S returns a relation containing all tuples in R or S  
*Intersection*: R ∩ S returns a relation containing all tuples in both R and S  
*Set-difference*: R - S returns a relation containing all tuples in R but not in S.

Union Compatability:

- Same number of attributes, and
- Corresponding attributes have the same domains
- No need to use the same attribute names

Cross-Product: ×

Given R(A, B, C) and S(X, Y), R × S = (A, B, C, X, Y), where every combination of A - Y is an entry in the new relation. ★ The set of attributes in R and S must be *disjoint*.

Inner Join

Include only tuples that satisfy the condition.

*θ*-join: R ⋈<sub>θ</sub> S

$$R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$$

Equi Join: R ⋈<sub>=</sub> S

Special type of *θ*-join where the only relational operator that can be used is equality.

$$R \bowtie_{=} S = \sigma_{=}(R \times S)$$

Natural Join: R ⋈ S

$$R \bowtie S = \pi_l(R \bowtie_c \rho_{a_1:b_1, \dots, a_n:b_n}(S))$$

where

- A = { a<sub>1</sub>, a<sub>2</sub>, ..., a<sub>n</sub> } is the set of common attributes between R and S
- c = (a<sub>1</sub> = b<sub>1</sub>) and ... and (a<sub>n</sub> = b<sub>n</sub>)
- l includes, in this order
  - List of attributes in R that are also in A
  - List of attributes in R that are not in A
  - List of attributes in S that are not in A

Outer Join

Include only tuples that do not satisfy the condition.

Dangling Tuple

Let attr(R) be the list of attributes in the schema of R. We say that t ∈ R is a dangling tuple in R wrt R ⋈ S if t ∉ π<sub>attr(R)</sub>(R ⋈<sub>c</sub> (S))

Left Outer Join: R ⋈<sub>[θ]</sub> S

Let null(R) denote a tuple of null values of same arity as R.

$$R \bowtie_{[\theta]} S = R \bowtie_{[\theta]} S \cup (dangle(R \bowtie_{[\theta]} S) \times \{null(S)\})$$

Right Outer Join: R ⋈<sub>[θ]</sub> S

$$R \bowtie_{[\theta]} S = R \bowtie_{[\theta]} S \cup (dangle(\{null(S)\} \times R \bowtie_{[\theta]} S))$$

Full Outer Join: R ⋈<sub>[θ]</sub> S

$$R \bowtie_{[\theta]} S = R \bowtie_{[\theta]} S \cup (dangle(R \bowtie_{[\theta]} S) \times \{null(S)\}) \cup (dangle(\{null(S)\} \times R \bowtie_{[\theta]} S))$$

Natural Left/Right/Outer Joins

Same schema as  $R \bowtie S$   
Joins

L	Operator	R	Visualization
Keep dangling			
		Keep dangling	
Keep dangling		Keep dangling	

Complex Equations - Equivalence

We say that 2 relational algebra expressions  $Q_1$  and  $Q_2$  are equivalent if for any input relations either: both produces error OR both produces the same result.  $\rightarrow$  **strongly** equivalent.  
★ **Weakly** equivalent can be defined as if there are no error then both produces the same result.

SQL - Structured Query Language

Domain-specific language designed for computations on relations.  
Made of: *Data definition, Data manipulation, Data Query, Data Control* and *Transaction Control*.

Data Types

- **boolean**: true/false (null == unknown)
- **integer**: signed 4 byte integer
- **float8**: double-precision floating point number (8 bytes)
- **numeric**: arbitrary precision floating point number
- **numeric(p, s)**: max total of p digits with max of s digits in fractional part
- **char(n)**: fixed-length string consisting of n characters
- **varchar(n)**: variable-length string up to n characters
- **text**: variable-length character string
- **date**: calender date (year, month, day)
- **timestamp**: date and time

Extended Data Types

- **Document**: XML/JSON
- **Spatial**: Point/Line/Polygon/Circle/Box/Path
- **Special**: Money/Currency & MAC/IP address

Create/Drop Table

```
1 CREATE TABLE Employees (  
2     id INTEGER,  
3     name VARCHAR(50),  
4     age INTEGER,  
5     role VARCHAR(50) DEFAULT 'sales'  
6 );
```

Insertion

```
1 INSERT INTO Employees  
2     VALUES (101, 'Sarah', 25, 'dev');  
3 INSERT INTO Employees (name, id)  
4     VALUES ('Judy', 102), ('Max', 103);
```

- Either **all** inserted or **none** inserted
- Attributes can be specied out-of-order (optionally)
- Missing values are replaced with NULL (if allowed) or default values (if specied)

Deleting

```
1 DELETE FROM <table_name>  
2     [ WHERE <condition> ];
```

- Condition is optional, unspecified: equivalent to always true
- Specified, condition can be arbitrarily complex

Integrity Constraints

**Principle of Acceptance:** Perform the operation if the condition **evaluates to TRUE**.  
**Principle of Rejection:** Reject the insertion if the condition **evaluates to FALSE**.

Type	Column	Table	Condition
Not-NULL	NOT NULL	-	IS NOT NULL
Unique	UNIQUE	UNIQUE(A <sub>1</sub> ,A <sub>2</sub> ,...)	x.A <sub>i</sub> <> y.A <sub>i</sub>
Primary Key	PRIMARY KEY	PRIMARY KEY(A <sub>1</sub> ,A <sub>2</sub> ,...)	UNIQUE & NOT NULL
Foreign Key	REFERENCES R <sub>1</sub> (B)	FOREIGN KEY (A <sub>1</sub> ,A <sub>2</sub> ,...) REFERENCES R <sub>1</sub> (B <sub>1</sub> ,B <sub>2</sub> ,...)	The tuple exists in R <sub>1</sub> or the tuple contains NULL value
General	CHECK (c)	CHECK (c)	Condition c does not evaluate to False

Is Null Predicate

x IS NULL: evaluates to true for null values, else false.  
x IS NOT NULL = NOT (x IS NULL)

Is Distinct from Predicate

x IS DISTINCT FROM y: equivalent to x <> y if both x and y are non-null else false if both are null, else true if only one is null  
x IS NOT DISTINCT FROM y = (x IS DISTINCT FROM y)

non-null Constraint

name        varchar(100) NOT NULL,

unique Constraint

studentId        INT UNIQUE, or  
unique (city, state) -- at bottom or  
primary key (sid, cid) -- at bottom

primary key Constraint

Primary key is a selected candidate key  $\rightarrow$  uniquely identifies a tuple in a relation and cannot be NULL, UNIQUE and NOT NULL  
studentId        INT CONSTRAINT stdnt\_pk PRIMARY KEY or  
PRIMARY KEY (eid, pname), -- at bottom

foreign key Constraint

studentId        INT REFERENCES Student (id) or FOREIGN KEY (a, b) REFERENCES Other (a, b) -- at bottom	
Keyword	Action
NO ACTION	<i>Reject</i> delete/update if it violates constraint (default value)
RESTRICT	Similar to "NO ACTION" except that check of constraint <i>cannot be deferred</i> (deferable constraints are discussed in a bit)
CASCADE	<i>Propagates</i> delete/update to the referencing tuples
SET DEFAULT	<i>Updates</i> the foreign key of the referencing tuples to some default value (important: default value must be a primary key in the referencing table)
SET NULL	<i>Updates</i> the foreign key of the referencing tuples to NULL value (important: corresponding column must be allowed to contain NULL values)

Modifying Database

Alter Table

Common changes: Adding/dropping columns, constraints or Changing specification of a column (data type, default values)

```
1 ALTER TABLE <table_name>  
2     [ALTER / ADD / DROP] [COLUMN / CONSTRAINT] <name>  
3     <changes>;
```

Drop Table

```
1 DROP TABLE  
2 [IF EXISTS] // no error if table does not exist  
3 <table_name>[, <table_name> [, <table_name> [...]]  
4 [CASCADE]; // also delete referencing table
```

Defferable Constraints

*unique, primary key* and *foreign key* constraints can be deferred using:

- DEFERRABLE INITIALLY DEFERRED: checked only at the end of the transaction.
- DEFERRABLE INITIALLY IMMEDIATE: checked after each statement.
- NOT DEFERRABLE: default for constraints

Considerations for deferrable constraints

- (+)

  - No need to care about the order of SQL statements within a transaction – allow intermediate state to temporarily violate constraints
  - Allows for cyclic foreign key constraints
  - Performance boost when constraint checks are bottlenecked – batch insert of larger number of tuples
- (-)

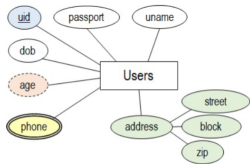
  - Troubleshooting can be more difficult
  - Data definition no longer unambiguous
  - May incur performance penalty when performing queries  $\rightarrow$  Certain checks may need to be done at run-time especially during the time when the constraint check is deferred

Entity-Relationship (ER) Model

*Entity*: Real-world object distinguishable from other objects  
*Attribute*: Specific information describing an entity, ovals  
*Entity set*: Collection of similar entities, rectangles  
*Key*: Represented as underlined attributes  
*Relationship*: Association among 2 or more entries  
*Relationship set*: Collection of similar relationships, represented by diamonds.  
Attributes used to describe information about relationships.

Subtypes:

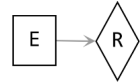
1. **Key Attributes** : *underlined*
  - Uniquely identifies each entity
2. **Composite Attributes** : *composed of other ovals*
  - Composed of multiple other attributes
3. **Multivalued Attributes** : *double-lined*
  - One or more values for a given entity
4. **Derived Attributes** : *dashed line*
  - Derived from other attributes



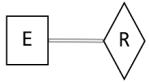
By default these relationships are *many-to-many*.

*Relationship role:* Shown explicitly when 1 entity appears 2 or more times in a relationship set  
*Degree:* An  $n$ -ary relationship set involves  $n$  entity roles; degree =  $n$ . ( $n = 2 \rightarrow$  binary,  $n = 3 \rightarrow$  ternary)  
*Relationship keys:* Each relationship set instance will have the primary keys of the entities as well as its own attributes. The primary keys of the relationship set will contain those primary keys as well as a subset of its own attributes, which will be underlined.

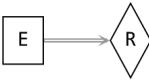
Relationship Constraints



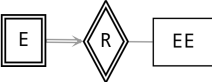
*Key Constraint:* Each instance of  $E$  can participate in **at most one instance** of  $R$ . Represented by an arrow. Allows for **one-to-many** if one entity has a constraint but the other does not or **one-to-one** if both entities have the constraint.



*Total Participation Constraint:* Each instance of  $E$  must participate in **at least one instance** of  $R$ . Represented by a double line. A single line is a *partial participation constraint* (i.e. 0 or more).



*Key & Total Participation Constraint:* Each instance of  $E$  participates in exactly one instance of  $R$ . Represented by double line arrow.



*Weak Entity Set:*  $E$  is a weak entity set with identifying owner  $EE$  & identifying relationship set  $R$ .  $E$  does not have its own key and requires the primary key of its owner entity to be uniquely identified. It must have a **many-to-one** relationship with  $EE$  **AND** total participation in  $R$ .

*Partial Key:* Set of attributes of a weak entity set that uniquely identifies a weak entity for a given owner entity.

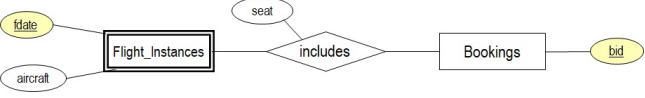
Relational Mapping

ER Diagram	Schema
Name of entity set	Name of table
Attribute of entity set	Attribute/column of table
Key attribute of entity set	Primary key of table
Derived attribute of entity set	<u>should not appear</u> *

When translating ERs to schemas, we can use the following questions as a guide:

- 1. Can **PK** be used to uniquely identify other attributes
- 2. Is the **LOWER BOUND** satisfied in the schema?
- 3. Is the **UPPER BOUND** satisfied in the schema?

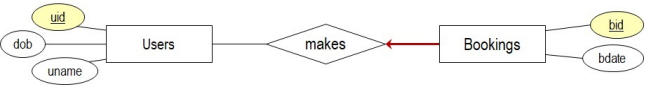
Many-to-Many



```
1 CREATE TABLE FlightInstances (  
2   fnum INT REFERENCES Flights,  
3   fdate DATE,  
4   aircraft VARCHAR(10),  
5   PRIMARY KEY (fnum, fdate)  
6 );  
7 CREATE TABLE Bookings (  
8   bid INT PRIMARY KEY  
9 );  
10 CREATE TABLE Includes (  
11   fdate DATE,  
12   fnum INT,  
13   bid INT REFERENCES Bookings,  
14   seat VARCHAR(10),  
15   REFERENCES FlightInstances,  
16   PRIMARY KEY (fnum, fdate, bid)  
17 );
```

- 1. Uniquely identify: (fnum, fdate, bid) able to ID the seat
- 2. Lower/Upper Bound: NA since it is a many-to-many

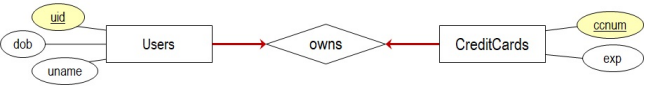
Many-to-One



```
1 CREATE TABLE Users (  
2   uid INT PRIMARY KEY,  
3   uname VARCHAR(100),  
4   dob DATE  
5 );  
6 -- Remove "Bookings", combine with "makes"  
7 CREATE TABLE MakesBookings (  
8   uid INT,  
9   bid INT,  
10  bdate DATE,  
11  FOREIGN KEY (uid) REFERENCES Users (uid),  
12  -- no more REFERENCES Bookings (bid)  
13  PRIMARY KEY (bid)  
14 );
```

- 1. Uniquely identify: (bid) able to ID the bdate
- 2. Lower Bound: Able to have a booking without user. **PK** is bid, user has no constraints (can be NULL).
- 3. Upper Bound: Upper limit is 1. Unable to have booking belonging to > 1 user (bid is **PK**).

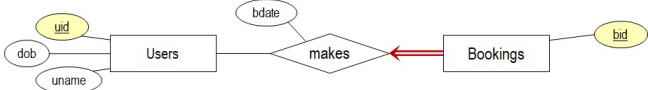
One-to-One



```
1 CREATE TABLE Users (  
2   uid INT PRIMARY KEY,  
3   ...  
4   REFERENCES CreditCards (ccnum)  
5 ); -- Assume CreditCards is not combined  
6 CREATE TABLE CreditCards (  
7   ccnum INT PRIMARY KEY,  
8   ...  
9   REFERENCES Users (uid)  
10 ); -- Assume Users is not combined
```

```
7   ccnum INT PRIMARY KEY,  
8   ...  
9   REFERENCES Users (uid)  
10 ); -- Assume Users is not combined
```

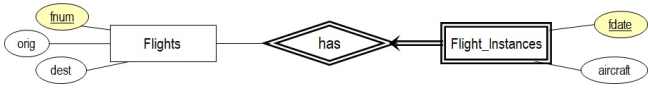
Key + Total Participation Constraint



```
1 CREATE TABLE Users (  
2   uid INT PRIMARY KEY,  
3   uname VARCHAR(100),  
4   dob DATE  
5 );  
6 CREATE TABLE MakesBookings (  
7   bid INT PRIMARY KEY,  
8   bdate DATE,  
9   uid INT NOT NULL,  
10  -- NOT NULL ensures that there must be a user!  
11  FOREIGN KEY (uid) REFERENCES Users (uid)  
12 );
```

- 1. Uniquely identify: (bid) able to ID the bdate
- 2. Lower Bound: Unable to have a booking without user. NOT NULL ensures each MakesBookings entry always has a user.
- 3. Upper Bound: Upper limit is 1. Unable to have booking belonging to > 1 user (bid is **PK**).

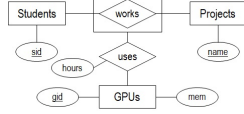
Weak Entity Sets



```
1 CREATE TABLE Users (  
2   fnum INT PRIMARY KEY,  
3   orig VARCHAR(10),  
4   dest VARCHAR(10)  
5 );  
6 CREATE TABLE FlightInstances (  
7   fnum INT REFERENCES Flights (fnum)  
8   ON DELETE CASCADE  
9   ON UPDATE CASCADE,  
10  fdate DATE,  
11  aircraft VARCHAR(10),  
12  PRIMARY KEY (fnum, fdate)  
13 );
```

Extended Notations - Aggregation

Abstraction that treats relationships as higher-level entities. Treat it as a relation class in OOP.

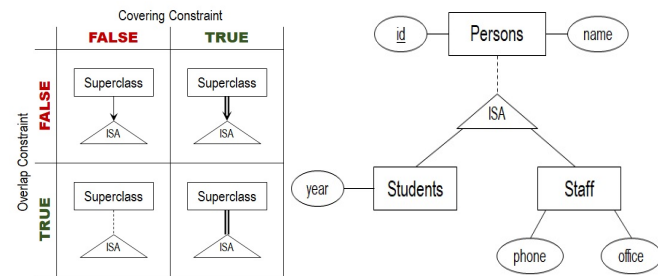


```
1 CREATE TABLE Uses (  
2   gid INT REFERENCES GPUs,  
3   sid INT,  
4   name VARCHAR(50),  
5   hours NUMERIC,  
6   PRIMARY KEY (gid, sid, pname),  
7   FOREIGN KEY (sid, pname)  
8     REFERENCES Works (sid, pname)  
9 );
```

## Extended Notations - ISA Hierarchies

Every entity subclass is an entity in its superclass.

**Overlap:** Can a superclass belong to multiple subclasses?  
**Covering:** Must a superclass belong to at least one subclass?



```
1 CREATE TABLE Persons (  
2   id   CHAR(20) PRIMARY KEY,  
3   name VARCHAR(50)  
4 );  
5 CREATE TABLE Students (  
6   id   CHAR(20) PRIMARY KEY,  
7   REFERENCES Persons (id) ON DELETE CASCADE  
8   year INT, -- or DATE?  
9 );  
10 CREATE TABLE Grads (  
11   id   CHAR(20) PRIMARY KEY  
12   REFERENCES Persons (id) ON DELETE CASCADE,  
13   phone INT,  
14   office VARCHAR(10)  
15 );
```

## SQL Queries

### SELECT-FROM-WHERE

```
1 SELECT [ DISTINCT ] <target-list>  
2 FROM <relation-list>  
3 WHERE <conditions>
```

SQL Query → RA Expression → Query Execution Plan → Executable Code  
EXAMPLE

```
1 SELECT DISTINCT a1, a2, ..., am  
2 FROM r1, r2, ..., rn  
3 WHERE c
```

RELATIONAL ALGEBRA:

$\pi[a_1, a_2, \dots, a_m](\sigma[c](r_1 \times r_2 \times \dots \times r_n))$

### Basic SQL Queries

- SELECT - Combine and process attribute values + Rename columns: Find restaurants, pizzas they sell and price of pizzas in SGD.

```
1 SELECT rname, pizza, '$$' || (price * 1.36) AS sgd  
2 FROM Sells;
```

- WHERE - Condition for NULL values: `x IS NULL` vs `x <> NOT NULL` vs `IS NULL`

SELECT cname	FROM Customers	WHERE area IS NULL;
Adi		
Yoga		

1 row

SELECT cname	FROM Customers	WHERE area = NULL;

0 rows

- WHERE - Pattern Matching

- matches any single character
- % matches any sequence of 0 or more characters
- EXAMPLE:

- \* Start with 'Ma', end with 'a':  
WHERE pizza LIKE 'Ma%a'
- \* starts with 'A' and consists of at least 5 characters:  
WHERE uname LIKE 'A\_\_\_\_%'

- SET OPERATIONS

```
1 Q1 UNION Q2  
2 Q1 INTERSECT Q2 -- treats everything as distinct  
3 Q1 EXCEPT Q2 -- treats everything as distinct  
4  
5 Q1 UNION ALL Q2 -- do not remove duplicates  
6 Q1 INTERSECT ALL Q2 -- do not remove duplicates  
7 Q1 EXCEPT ALL Q2 -- do not remove duplicates
```

- JOIN OPERATIONS: Most common in practice: cross product + selection condition + attribute selection

```
1 SELECT DISTINCT rname, price  
2 FROM Sells S JOIN Recipes R  
3   ON S.pizza = R.pizza  
4 WHERE R.ingredient = 'Cheese';  
5  
6 SELECT DISTINCT rname, price  
7 FROM Sells S NATURAL JOIN Recipes R  
8 -- NATURAL JOIN Matches all common attributes  
9 WHERE R.ingredient = 'Cheese';  
10  
11 SELECT DISTINCT C.cname  
12 FROM Customers C LEFT JOIN Likes L  
13   ON C.cname = L.cname  
14 WHERE L.pizza IS NULL; -- keep dangling tuples
```

### Composing

#### Scalar Subqueries:

A query that returns at most a single value.

```
1 SELECT *  
2 FROM Sells  
3 WHERE price < (SELECT S.price FROM Sells S  
4               WHERE rname = 'Corleone Corner');
```

#### [ NOT ] IN:

Subquery **MUST** return exactly one column.

- IN returns TRUE if <expr> matches any subquery row
- NOT IN returns TRUE if <expr> matches no subquery row

```
1 WHERE <expr> [ NOT ] IN [ <subquery> | <tuple> ];
```

#### ANY/SOME

Subquery **MUST** return exactly one column. Expression <expr> is compared to each row from subquery using the operator <op>.

- ANY returns TRUE if comparison evaluates to TRUE for **at least one** row in the subquery

```
1 WHERE <expr> <op> ANY <subquery>;
```

## ALL

Subquery **MUST** return exactly one column. Expression <expr> is compared to each row from subquery using the operator <op>.

- ALL returns TRUE if comparison evaluates to TRUE for **ALL** rows in the subquery

```
1 WHERE <expr> <op> ALL <subquery>;
```

#### [ NOT ] EXISTS

May return any number of columns

- EXISTS returns TRUE if the subquery returns at least one row
- NOT EXISTS returns TRUE if the subquery returns no row

```
1 WHERE [ NOT ] EXISTS <subquery>;
```

### Scoping

- A table alias declared in a (sub-)query Q can only be used in Q or subqueries nested within Q
- If the same table alias is declared in both subquery Q<sub>1</sub> and in outer query Q<sub>0</sub> (or not at all), the declaration in Q<sub>1</sub> is applied

```
1 SELECT DISTINCT S1.rname  
2 FROM Sells S1  
3 WHERE price < ANY -- S1.price  
4   (SELECT price -- S2.price  
5    FROM Sells S2  
6    WHERE S2.rname = 'Lorenzo Tavern');
```

\* Not all constructs required: IN  $\equiv$  ANY  $\equiv$  EXIST

### Ordering: ORDER BY

- ORDER BY <attribute> ASC: Ascending Order (default for SQL, ASC can be removed)
- ORDER BY <attribute> DESC: Descending Order
- If duplicate removal needed, attribute being sorted **must** appear in SELECT
- Sorting w.r.t multiple attributes / differing order supported

SQL	Result																		
<pre>SELECT * FROM Recipes WHERE pizza IN (SELECT pizza FROM Recipes WHERE ingredient = 'Cheese') ORDER BY pizza DESC, ingredient ASC;</pre>	<table><thead><tr><th>pizza</th><th>ingredient</th></tr></thead><tbody><tr><td>Siciliana</td><td>Anchovies</td></tr><tr><td>Siciliana</td><td>Caspers</td></tr><tr><td>Siciliana</td><td>Cheese</td></tr><tr><td>Margherita</td><td>Cheese</td></tr><tr><td>Margherita</td><td>Tomato</td></tr><tr><td>Diavola</td><td>Cheese</td></tr><tr><td>Diavola</td><td>Chilli</td></tr><tr><td>Diavola</td><td>Salami</td></tr></tbody></table>	pizza	ingredient	Siciliana	Anchovies	Siciliana	Caspers	Siciliana	Cheese	Margherita	Cheese	Margherita	Tomato	Diavola	Cheese	Diavola	Chilli	Diavola	Salami
pizza	ingredient																		
Siciliana	Anchovies																		
Siciliana	Caspers																		
Siciliana	Cheese																		
Margherita	Cheese																		
Margherita	Tomato																		
Diavola	Cheese																		
Diavola	Chilli																		
Diavola	Salami																		

### Ordering: LIMIT and OFFSET

- LIMIT k Return the first\* k rows of the result table
- OFFSET i Specify the position of the first row to be considered
- \* LIMIT and OFFSET are typically meaningful only in combination with ORDER BY

Find the next 3 cheapest pizza, restaurant selling them and price.

```
1 SELECT *  
2 FROM Sells  
3 ORDER BY price ASC  
4 OFFSET 3  
5 LIMIT 3;
```



Aggregation

Function	Input Type	Output Type
MIN	any comparable type	same as input
MAX	any comparable type	same as input
SUM	Numeric data (e.g., INT, BIGINT, REAL, etc)	SUM(INT) → BIGINT; SUM(REAL) → REAL
COUNT	any data	BIGINT

Aggregate functions compute a *single value* from a set of tuples.

```
1 SELECT MIN(price) AS lowest
2 FROM Sells;
```

Query	Interpretation	Result
SELECT MIN(A) FROM R;	Minimum <i>non-NULL</i> value in A	0
SELECT MAX(A) FROM R;	Maximum <i>non-NULL</i> value in A	42
SELECT AVG(A) FROM R;	Average of <i>non-NULL</i> values in A	12 (i.e., 48 / 4)
SELECT SUM(A) FROM R;	Sum of <i>non-NULL</i> values in A	48
SELECT COUNT(A) FROM R;	Count of <i>non-NULL</i> values in A	4
SELECT COUNT(*) FROM R;	Count of rows in A <b>(all rows)</b>	5
SELECT AVG(DISTINCT A) FROM R;	Average of <i>distinct</i> <i>non-NULL</i> value in A	15
SELECT SUM(DISTINCT A) FROM R;	Sum of <i>distinct</i> <i>non-NULL</i> value in A	45
SELECT COUNT(DISTINCT A) FROM R;	Count of <i>distinct</i> <i>non-NULL</i> value in A	3

Left: R is an empty relation with attribute B  
Right: S is a non-empty relation with n-rows with attribute A only having NULL values

Query	Result	Query	Result
SELECT MIN(B) FROM R;	NULL	SELECT MIN(A) FROM S;	NULL
SELECT MAX(B) FROM R;	NULL	SELECT MAX(A) FROM S;	NULL
SELECT AVG(B) FROM R;	NULL	SELECT AVG(A) FROM S;	NULL
SELECT SUM(B) FROM R;	NULL	SELECT SUM(A) FROM S;	NULL
SELECT COUNT(B) FROM R;	0	SELECT COUNT(A) FROM S;	0
SELECT COUNT(*) FROM R;	0	SELECT COUNT(*) FROM S;	n

Grouping

Application of aggregate functions are over all tuples of a relation

- Logical partition of relation into groups based on values for specific attributes
- Application of aggregation functions are now over each group (one result tuple for each group)

```
1 GROUP BY attr1, attr2, ...
```

★ f column $A_i$ appears in SELECT, one condition must hold
• $A_i$ appears in the GROUP BY clause
• $A_i$ appears as input of aggregation function in SELECT clause
• The pkey of R appears in the GROUP BY clause

EXAMPLE

The first query is valid because it groups the results by the unique **rname** column from the **Restaurants** table and counts the number of **cname** values from the **Customers** table that have the same area value as the grouped **rname**. This is a valid aggregation because the **GROUP BY** clause includes all non-aggregated columns (i.e. **rname** and **area**) from the **Restaurants** table and the aggregated column (i.e. **cname**) from the **Customers** table.

The second query is not valid because it groups the results by the unique **rname** column from the **Restaurants** table and also includes the non-aggregated **cname** column from the **Customers** table in the **SELECT** clause. When using **GROUP BY**, all non-aggregated columns in the **SELECT** clause must also be included in the **GROUP BY** clause. In this case, the **cname**

column from the **Customers** table is not included in the **GROUP BY** clause, so it is not clear which **cname** value should be displayed for each group of **rname** values.

```
SELECT R.rname, R.area,
       COUNT (C.cname)
FROM   Restaurants R, Customers C
WHERE  R.area = C.area
GROUP BY R.rname;
```

```
SELECT R.rname, C.cname,
       COUNT (C.cname)
FROM   Restaurants R, Customers C
WHERE  R.area = C.area
GROUP BY R.rname;
```

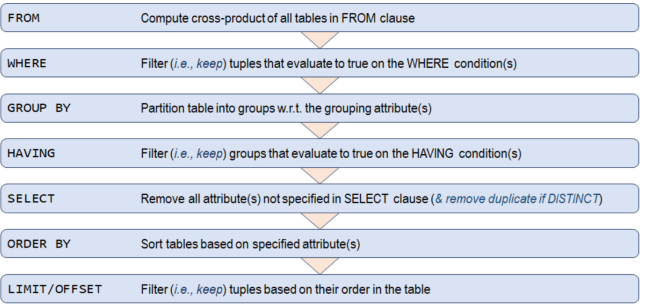
Having

**WHERE** clauses only works on each row. If we were to use **WHERE p < AVG(p)**, any p that does not fit this condition will be removed, resulting in **AVG(p)** to change. **HAVING** instead aggregates the condition for each group defined by **GROUP BY**.

```
1 GROUP BY attr1, attr2, ...
2 HAVING <condition>
```

★ If column $A_i$ appears in HAVING, one condition must hold
• $A_i$ appears in the GROUP BY clause
• $A_i$ appears as input of aggregation function in HAVING clause
• The pkey of R appears in the GROUP BY clause

Conceptual Evaluation of Queries



Conditional Expressions: CASE

Only one of the results will be returned, else is optional, and NULL returned if no condition matched.

```
1 CASE
2   WHEN <condition_1> THEN <result_1>
3   WHEN <condition_2> THEN <result_2>
4   ELSE result_0
5 END
```

Conditional Expressions: COALESCE

Returns the first non-NULL value in the list of input arguments (order of **order** matters!)

```
1 COALESCE(<value_1>, <value_2>, <value_3>, ...)
2 // the above looks at <value_1> first, if it is
3 // non-NULL, it will return <value_1>
```

Conditional Expressions: NULLIF

Returns NULL if **<value\_1> = <value\_2>**, otherwise returns **<value\_1>**

```
1 NULLIF(<value_1>, <value_2>)
```

<pre>SELECT   NULLIF(1,1) AS val;</pre>	<pre>SELECT   NULLIF(2,1) AS val;</pre>	<pre>SELECT   NULLIF(NULL,1) AS val;</pre>	<pre>SELECT   NULLIF(2,NULL) AS val;</pre>
<pre>val NULL</pre>	<pre>val 2</pre>	<pre>val NULL</pre>	<pre>val 2</pre>

Common Table Expression

- CTE<sub>i</sub> is the name of a temporary table defined by query Q<sub>i</sub>
- Each CTE<sub>i</sub> can reference any other CTE<sub>j</sub> that has been declared before CTE<sub>i</sub> (i.e., j < i)
- The main SELECT statement Q<sub>0</sub> can reference any possible subset of all CTE<sub>i</sub> → Any CTE<sub>i</sub> not referenced can be deleted

```
1 WITH
2   CTE_1 AS ( Q_1 ),
3   CTE_2 AS ( Q_2 ),
4   ...
5   CTE_n AS ( Q_n )
6 Q_0
7 -- main SELECT statement
```

Views

Used since we only usually need parts of the table, restrict access to table for certain users, and often use the same queries and subqueries frequently.

```
1 CREATE VIEW <name> AS
2 <query> -- SELECT statement
3 ;
```

- A VIEW is a permanently named query, computation done each time the virtual table is accessed
- Results are not permanently stored, no actual table created.

Extended Concepts

Universal Quantification

QUERY: restaurant that sells all pizzas liked by ‘Homer’

Transformation:

→ there does not exist pizza that ‘Homer’ Likes and not sold by the restaurant: all pizzas that ‘Homer’ likes and restaurant with **rname = R** does not sell

```
1 SELECT L.pizza
2 FROM Likes L
3 WHERE L.cname = 'Homer'
4       AND NOT EXISTS (
5         SELECT 1
6         FROM Sells S
7         WHERE S.pizza = L.pizza
8         AND S.rname = 'R'
9       );
```

→ Find no possible restaurants that does not sell any pizza liked by ‘Homer’.

```
1 SELECT DISTINCT S1.rname
2 FROM Sells S1
3 WHERE NOT EXISTS (
4   SELECT 1 FROM Likes L
5   WHERE L.cname = 'Homer'
6         AND NOT EXISTS (
7           SELECT 1 FROM Sells S
8           WHERE S.pizza = L.pizza
```

```

9      AND S.rname = S1.rname
10    )
11  );

```

### Cardinality

→ Find the set of pizzas sold by a restaurant that is a SUPERSET or EQUAL to that of the pizzas liked by ‘Homer’.

```

1 SELECT DISTINCT rname
2 FROM Sells S
3 WHERE (
4   SELECT COUNT(DISTINCT pizza)
5   FROM (
6     SELECT pizza FROM Sells S1 WHERE S1.rname = S.←
7     rname      // selecting correct R
8   UNION
9   SELECT pizza FROM Likes WHERE cname = 'Homer' ←
10  ) AS T1
11 ) = (
12   SELECT COUNT(DISTINCT pizza)
13   FROM Sells S1 WHERE S1.rname = S.rname      // ←
14   counting the number of pizzas
15 );

```

### Recursive Queries

```

1 WITH RECURSIVE
2   CTE_name AS (
3     Q_1
4     UNION [ ALL ]
5     Q_2 ( CTE_name )
6   )
7 Q_0 ( CTE_name )

```

- Q<sub>1</sub> is non-recursive
- Q<sub>2</sub> is recursive and can reference CTE<sub>name</sub>
- Query is evaluated ‘lazily’, stops when a fixed-point is reached

Find all MRT stations can be reached from NS1 in at most 3 stops

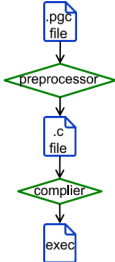
```

1 WITH RECURSIVE
2   Linker(to_stn, stops) AS (
3     SELECT to_stn, 0
4     FROM MRT
5     WHERE fr_stn = 'NS1'
6   UNION ALL
7   SELECT M.to_stn, L.stops + 1
8   FROM Linker L, MRT M
9   WHERE L.to_stn = M.fr_stn
10   AND L.stops < 2
11 )
12 SELECT DISTINCT (to_stn)
13 FROM Linker;

```

### Dynamic SQL

If we need to write different programs to access a database, we can mix a host language with SQL. The program is passed into a preprocessor, which returns a program in the host language. After which, this program is compiled into executable code.



### Static SQL

```

1 int main() {
2   EXEC SQL BEGIN DECLARE SECTION;      // 1.
3   char supplName[30], prodName[20]; float price;
4   EXEC SQL END DECLARE SECTION;
5
6   EXEC SQL CONNECT TO @localhost USER john; // 2.
7
8   // some code that assigns values to // 3.
9   // supplName, prodName, price;
10
11  EXEC SQL INSERT INTO
12    Sells(supplName, prodName, price) // 4.
13    VALUES(:supplName, :prodName, :price);
14  EXEC SQL DISCONNECT; // 5.
15  return 0;
16 }

```

1. Declares **shared** variables that can be used by both the host language and SQL statements
2. Specifies the database to connect to and ther username to use
3. Host part of the language
4. Inserts record into the table, must be preceded by a colon → this is referred to as a static SQL
5. Closes the database connection

### Dynamic SQL

The SQL statement to be executed is not fixed in advance. Instead, we store the statement in a string variable, which is then compiled and executed at runtime, checking the syntax at runtime.

- EXEC SQL EXECUTE IMMEDIATE :stmt, this part executes the SQL statement stored in stmt.

```

1 int main() {
2   EXEC SQL BEGIN DECLARE SECTION;      // 1.
3   const char *stmt = "INSERT INTO test VALUES←
4     (?, ?);";
5   EXEC SQL END DECLARE SECTION;
6
7   EXEC SQL CONNECT TO @localhost USER john; // 2.
8
9   EXEC SQL PREPARE mystmt FROM :stmt; // 3.
10
11  EXEC SQL EXECUTE mystmt USING 42, 'foobar'; // 4.
12
13  EXEC SQL DEALLOCATE PREPARE mystmt; // 5.
14
15  EXEC SQL DISCONNECT;
16  return 0;
17 }

```

1. Init a SQL statement with 2 parameters
2. Specifies the database to connect to and ther username to use
3. Parsing the SQL statement and gives it a name mystmt
4. Executes SQL statement using 42 and 'foobar' as parameters
5. Releases the resources allocated to mystmt (EVERY PREPARE statement should be deallocated when no longer needed)

### Call-Level Interface (CLI)

e.g. Java Database Connectivity (JDBC) or Open Database Connectivity (ODBC). This is used when we do not want to mix the host languages with SQL. Instead of using EXEC SQL, we instead call an API to run the queries.

### ODBC Example in C++

```

1 void main() {
2   char *sql;
3   connection C("dbname= testbd user=postgre \ ←
4     password=cohondob hostaddr=127.0.0.1 port←
5     =5432"); // 1.
6
7   // Some code here
8
9   // Create SQL statement
10  sql = "CREATE TABLE COMPANY (" \
11    "ID INT PRIMARY KEY NOT NULL," \
12    "NAME TEXT NOT NULL);"; // 2.
13
14  work W(C); // Create a transactional object 3.
15
16  W.exec( sql ); // 4.
17  W.commit();
18  C.disconnect(); // 5.
19 }

```

1. Creates a connection to the database
2. Assigns the SQL statement as a string to the sql variable
3. creates an object W for executing the SQL statement
4. Executes SQL statement recorded in sql
5. Closes database connection

### Functions

Makes code reusable, easy to maintain, performance and acheives security (more in SQL injections). ISO standard: SQL/PSM (Persistant Stored Modules). PostgreSQL: PL/pgSQL

#### Example 1 — Convert marks to letter grades

- [70,100] → A
- [50,60] → C
- [60,70] → B
- [0,50] → D

```

1 CREATE OR REPLACE FUNCTION convert(mark INT)
2 RETURNS char(1) AS $$
3   SELECT CASE
4     WHEN mark >= 70 THEN 'A'
5     WHEN mark >= 60 THEN 'B'
6     WHEN mark >= 50 THEN 'C'
7     ELSE 'D'
8   END;
9 $$ LANGUAGE sql;
10
11 /*
12 Query:SELECT convert (90);
13 Result: A
14
15 Query: SELECT Name, convert( Mark ) FROM Scores;
16
17 +-----+-----+
18 | Name | Mark |
19 +-----+-----+
20 | Alice | 92 |
21 | Bob | 63 |
22 | Cathy | 58 |
23 | David | 47 |
24 +-----+-----+
25 */

```

Name	Mark	Convert
Alice	92	A
Bob	63	B
Cathy	58	C
David	47	D

## Example 2 — Return all tuples with highest mark

```

1 CREATE OR REPLACE FUNCTION topStudents()
2 RETURNS SETOF Scores AS $$
3     SELECT *
4     FROM Scores
5     WHERE MARK = (SELECT MAX(Mark) FROM Scores);
6 $$ LANGUAGE sql;
7
8 /*
9 If we only want 1 student, we can instead do:
10 ORDER BY Mark DESC LIMIT 1 to find the TOP student.
11
12 Query: SELECT * FROM topStudents();
13
14 +-----+-----+
15 | Name | Mark |
16 +-----+-----+
17 | Alice | 92 |
18 | Bob | 63 |
19 | Cathy | 58 |
20 | David | 92 |
21 +-----+-----+
22 */

```

Name	Mark	Top Mark	Cnt
Alice	92	92	2
Bob	63	63	1
Cathy	58	58	1
David	92		

### Difference between SETOF, RECORDS, TABLE

- SETOF is used when we want to return multiple records
- RECORDS is used for only 1 record, and it is not in the current schema
- SETOF RECORDS is used for multiple tuples, which are not in the current schema
- TABLE is similar to SETOF RECORDS, this means our earlier code in Example 4 can also be done with TABLE

```

1 CREATE OR REPLACE FUNCTION markCnt()
2 RETURNS TABLE(Mark INT, Cnt INT) AS $$
3     SELECT Mark, COUNT(*)
4     FROM Scores
5     GROUP BY Mark;
6 $$ LANGUAGE sql;

```

## Example 3 — Returning a new tuple

```

1 CREATE OR REPLACE FUNCTION topMarkCnt(OUT TopMark INT↵
2 , OUT Cnt INT)
3 RETURNS RECORD AS $$
4     SELECT Mark, COUNT(*)
5     FROM Scores
6     WHERE MARK = (SELECT MAX(Mark) FROM Scores)
7     GROUP BY Mark;
8 $$ LANGUAGE sql;
9
10 /*
11 This returns the highest mark and its number of ↵
12 occurrences
13
14 Query: SELECT topMarkCnt();
15
16 +-----+-----+
17 | Name | Mark |
18 +-----+-----+
19 | Alice | 92 |
20 | Bob | 63 |
21 | Cathy | 58 |
22 | David | 92 |
23 +-----+-----+
24 */

```

### Procedures

If no return value/tuple is needed we may use **SQL procedures** instead. To execute the following procedure, we use:  
CALL transfer('Alice', 'Bob', 100);

### Procedure Example

```

1 CREATE OR REPLACE PROCEDURE transfer(fromAcc TEXT, ↵
2 toAcct TEXT, amount INT)
3 AS $$
4     UPDATE Accounts
5     SET balance = balance - amount
6     WHERE name = fromAcct;
7
8     UPDATE Accounts
9     SET balance = balance + amount
10    WHERE name = toAcct;
11 $$ LANGUAGE SQL

```

SQL Functions/Procedures can be complex, having variables and *control structures*

- IF ... THEN ... ELSE ... END IF
- LOOP ... END LOOP
- EXIT WHEN – Used within a loop to exit the loop
- WHILE ... LOOP ... END LOOP
- FOR ... IN ... LOOP ... END LOOP

### IF THEN Example

```

1 CREATE OR REPLACE FUNCTION swap(INOUT val1 INT, INOUT↵
2 val2 INT)
3 RETURNS RECORD AS $$
4 DECLARE
5     temp_val INTEGER;
6 BEGIN
7     IF val1 > val2 THEN
8         temp_val := val1;
9         val1 := val2;
10        val2 := temp_val;
11    END IF;

```

```

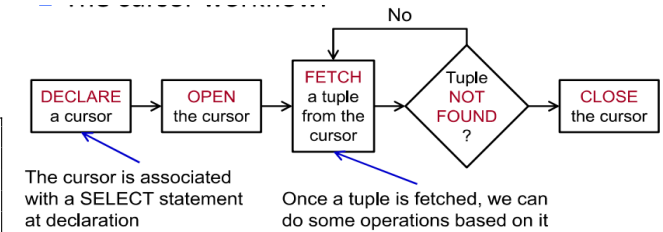
11 END;
12 $$ LANGUAGE plpgsql;

```

- SELECT swap(99, 11); → (11, 99)
- SELECT swap(11, 99); → (11, 99)

### Cursors

Cursors are used to loop over a sorted sequence within the schema. If we need to loop over them and do some computation, we can use a cursor as a pointer to reference every variable in the table iteratively. A cursor enables us to access each individual row returned by a SELECT statement.



```

1 CREATE OR REPLACE FUNCTION score_gap()
2 RETURNS TABLE(name TEXT, mark INT, gap INT) AS $$
3 DECLARE // 1.
4     curs CURSOR FOR (SELECT * FROM Scores ORDER BY ↵
5         Mark DESC);
6     r RECORD;
7     prv_mark INT;
8 BEGIN
9     prv_mark := -1;
10    OPEN curs; // 2.
11    LOOP
12        FETCH curs INTO r; // 3.
13        EXIT WHEN NOT FOUND; // 4.
14        name := r.Name; // 5.
15        mark := r.Mark;
16        IF prv_mark >= 0 THEN gap := prv_mark - mark;
17        ELSE gap := NULL;
18        END IF;
19        RETURN NEXT; // 6.
20        prv_mark := r.Mark; // 7.
21    END LOOP;
22    CLOSE curs; // 8.
23 END;
24 $$ LANGUAGE plpgsql;

```

1. Declares cursor variable
2. Executes SQL statement, lets curs point to beginning of result
3. Read next tuple from curs and put into r
4. If FETCH operation did not get any tuple, terminate loop
5. If FETCH operation got tuple, compute values of name, mark and gap
6. Inserts a tuple (name, mark, gap) to the output of function
7. Updates prv\_mark variable
8. Release resources allocated to curs

### Other Cursor Movements

- FETCH PRIOR FROM cur INTO r → fetch the prior row
- FETCH FIRST FROM cur INTO r → fetch first row
- FETCH LAST FROM cur INTO r → fetch last row
- FETCH ABSOLUTE 3 FROM cur INTO r → fetch the 3<sup>rd</sup> tuple

## Example 4 — Return new set of tuples

```

1 CREATE OR REPLACE FUNCTION markCnt(OUT TopMark INT, ↵
2 OUT Cnt INT)
3 RETURNS SETOF RECORD Scores AS $$
4     SELECT Mark, COUNT(*)
5     FROM Scores
6     GROUP BY Mark;
7 $$ LANGUAGE sql;
8
9 /*
10 Returns distinct Mark and number of occurrences
11
12 Query: SELECT markCnt();
13
14 +-----+-----+
15 | Name | Mark |
16 +-----+-----+
17 | Alice | 92 |
18 | Bob | 63 |
19 | Cathy | 58 |
20 | David | 92 |
21 +-----+-----+
22 */

```

## SQL Injections

Type of attack on dynamic SQL. Imagine `stmt` is the concatenation of 3 sub-strings:

```
SELECT * FROM T WHERE Name = ' input_name ';
```

Someone with malicious intent can instead put

`a'`; DROP TABLE T; SELECT 'a into *input\_name*. Now, T will be deleted once `stmt` is executed.

## Triggers ↔ Trigger Functions

We want to implement this on an insertion, need the database to check this particular **condition** whenever appropriate.

- Expression condition about 'insertion occurring on Scores' → Trigger Functions
- Database to check this condition whenever appropriate → Triggers

```
1 // Trigger
2 CREATE TRIGGER scores_log_trigger
3 AFTER INSERT ON Scores
4 FOR EACH ROW EXECUTE FUNCTION scores_log_function();
5
6 // Trigger Function
7 CREATE OR REPLACE FUNCTION scores_log_function() ↔
8 RETURNS TRIGGER AS $$
9 BEGIN
10     INSERT INTO Scores_Log(Name, EntryDate)
11     VALUES (NEW.Name, CURRENT_DATE);
12     RETURN NULL;
13 END;
14 $$ LANGUAGE plpgsql;
```

- Trigger tells database to watch out for insertions on Scores and call the `scores_log_func()` after each insertion of a tuple
- 'RETURNS TRIGGER' indicates that this is a trigger function, only RETURNS TRIGGER allowed as only TRIGGER has access to the NEW.
- NEW refers to new row inserted into Scores
- CURRENT\_DATE returns the current date

### What else can a trigger function access?

- TG\_OP: operation that activates trigger: INSERT, UPDATE, DELETE
- TG\_TABLE\_NAME: name of table that cause trigger invocation
- OLD: old tuple being updated/deleted

### Trigger Timing Example

Whenever there is a insert/delete/update, insert tuple into `Scores_Log2` to record: name, operation performed, date of operation

```
1 // Trigger
2 CREATE TRIGGER scores_log_trigger2
3 AFTER INSERT OR DELETE OR UPDATE ON Scores
4 FOR EACH ROW EXECUTE FUNCTION scores_log2_func();
5
6 // Trigger Function
7 CREATE OR REPLACE FUNCTION scores_log2_function()
8 RETURNS TRIGGER AS $$ BEGIN
9     IF (TG_OP = 'INSERT') THEN
10         INSERT INTO Scores_log2
11         SELECT NEW.Name, 'Insert', CURRENT_DATE;
12         RETURN NEW;
13     ELSEIF (TG_OP = 'DELETE') THEN
14         INSERT INTO Scores_log2
15         SELECT OLD.Name, 'Delete', CURRENT_DATE;
```

```
16     RETURN OLD;
17     ELSEIF (TG_OP = 'UPDATE') THEN
18         INSERT INTO Scores_log2
19         SELECT NEW.Name, 'Update', CURRENT_DATE;
20     RETURN NEW;
21 END IF;
22 END;
23 $$ LANGUAGE plpgsql;
```

- AFTER: indicates that `scores_log2_func()` is executed after the insertion on Scores is done
- BEFORE: `scores_log2_func()` would be executed **before** the insertion
- INSTEAD OF: `scores_log2_func()` would be executed **instead of** insertion

### BEFORE Trigger Example

```
1 // Trigger
2 CREATE TRIGGER for_Elise_trigger
3 BEFORE INSERT ON Scores
4 FOR EACH ROW EXECUTE FUNCTION for_Elise_func();
5
6 // Trigger Function
7 CREATE OR REPLACE FUNCTION for_Elise_func()
8 RETURNS TRIGGER AS $$ BEGIN
9     IF (NEW.Name = 'Elise') THEN
10         NEW.Mark := 100
11     ENDIF;
12     RETURN -----;
13 END;
14 $$ LANGUAGE plpgsql;
```

1. NEW: Elise's mark would be 100 regardless of what we insert.
2. NULL: Signals to postgres that this execution has a problem, no tuple can be inserted. RETURN NULL in a BEFORE trigger tells database to ignore rest of operation.
3. OLD: No tuple can be inserted, for INSERT, OLD is initially set to NULL, thus RETURN OLD is same as RETURN NULL.
4. OLD exception: Whenever the function returns a **non-null** tuple, trigger would use it as a tuple to be inserted

```
1 CREATE OR REPLACE FUNCTION for_Elise_func()
2 RETURNS TRIGGER AS $$
3 BEGIN
4     OLD.Name := 'Haha';
5     OLD.Mark := 0;
6     RETURN OLD;
7 END;
8 $$ LANGUAGE plpgsql;
```

### Return Values of Trigger Function — BEFORE

- BEFORE INSERT:
  - Returning a non-null tuple `t`, `t` will be inserted
  - Returning a null tuple, no tuple inserted

BEFORE UPDATE:

- Returning a non-null tuple `t`, `t` will be updated tuple
- Returning a null tuple, no tuple updated

BEFORE DELETE:

- Returning a non-null tuple `t`, deletion proceeds as normal
- Returning a null tuple, no tuple deleted

### Return Values of Trigger Function — AFTER

- AFTER INSERT, AFTER UPDATE, AFTER DELETE, does not matter what is returned, as trigger function is invoked **after** main operation is done.

### INSTEAD OF Trigger Example

This type of triggers can be defined as VIEWS only. Instead of doing something on a view, do it on a table. e.g. Whenever someone wants to update tuple in Max\_Score (VIEW), update corresponding tuple in Scores.

```
1 // Trigger
2 CREATE TRIGGER update_max_trigger
3 INSTEAD OF UPDATE ON Max_Score
4 FOR EACH ROW EXECUTE FUNCTION update_max_func();
5
6 // Trigger Function
7 CREATE OR REPLACE FUNCTION update_max_func()
8 RETURNS TRIGGER AS $$
9 BEGIN
10     UPDATE Scores
11     SET Mark = NEW.Mark WHERE Name = OLD.Name;
12     RETURN NEW;
13 END;
14 $$ LANGUAGE plpgsql;
```

- Returning NULL signals the database to ignore the rest of the operation on the current row.
- Returning a non-null tuple signals the database to proceed as normal.

### Trigger Levels

FOR EACH ROW: **row-level** trigger that executes trigger function for every tuple encountered.

FOR EACH STATEMENT: **statement-level** trigger that executes the trigger function only once

```
1 // Trigger
2 CREATE TRIGGER del_warn_trigger
3 BEFORE DELETE ON Scores_Log
4 FOR EACH STATEMENT EXECUTE FUNCTION del_warn_func();
5
6 // Trigger Function
7 CREATE OR REPLACE FUNCTION del_warn_trigger()
8 RETURNS TRIGGER AS $$
9 BEGIN
10     RAISE NOTICE 'Not supposed to delete from log.';
11     RETURN NULL; /* db prompts on deletion attempt */
12 END;
13 $$ LANGUAGE plpgsql;
```

- Statement-level triggers ignores the values returned by trigger functions.
- RETURN NULL would not make database omit subsequent operations

### EXCEPTION Example

```
1 // Same Trigger as above
2 // Trigger Function
3 CREATE OR REPLACE FUNCTION del_warn_trigger()
4 RETURNS TRIGGER AS $$
5 BEGIN
6     RAISE EXCEPTION 'No deletion from log allowed.';
7     RETURN NULL;
8 END;
9 $$ LANGUAGE plpgsql;
```



- INSTEAD OF only allowed on **row-level**
- BEFORE/AFTER allowed on both **row** and **statement-level**

### Trigger Condition

```
1 // Trigger
2 CREATE TRIGGER for_Elise_trigger
3 BEFORE INSERT ON Scores
4 FOR EACH ROW WHEN (NEW.Name = 'Elise')
5 EXECUTE FUNCTION for_Elise_func();
6
7 // Trigger Function
8 CREATE OR REPLACE FUNCTION for_Elise_func()
9 RETURNS TRIGGER AS $$
10 BEGIN
11     NEW.Mark := 100;
12     RETURN NEW;
13 END;
14 $$ LANGUAGE plpgsql;
```

In general, the condition in WHEN() could be more complicated, subject to the following requirements:

- No SELECT in WHEN()
- No NEW in WHEN() for DELETE
- No OLD in WHEN() for INSERT
- No WHEN for INSTEAD OF

### Deffered Trigger

Scenarios where we need to **defer** checking of triggers. e.g. trigger on INSERT/UPDATE/DELETE checking total balance of each customer Accounts. Total balance should be at least **150**.

- Alice transfer 100 from Account 1 to Account 2, assuming both have a starting balance of 100. 2 update statements, deduct 100 from Account 1, add 100 to Account 2
- Trigger requirement violated in first update statement
- **Solution:** Put the 2 update statements in 1 transaction, defer the trigger check to the end of the transaction

```
1 CREATE CONSTRAINT TRIGGER bal_check_trigger
2 AFTER INSERT OR UPDATE OR DELETE ON Account
3 DEFERRABLE INITIALLY DEFERRED
4 FOR EACH ROW EXECUTE FUNCTION bal_check_func();
```

- CONSTRAINT and DEFERRABLE together indicate that the trigger can be deferred
- INITIALLY DEFERRED indicates that by default, trigger is deferred. INITIALLY IMMEDIATE trigger not deferred by default.

DEFERRED Triggers only work for AFTER and FOR EACH ROW

- AFTER to defer the trigger, has to be allowed to execute after main operation

With the above deffered trigger, we can do the following, as the trigger will only be activated at COMMIT

```
1 BEGIN TRANSACTION;
2 UPDATE ACCOUNT SET Bal = Bal - 100 WHERE AID = 1;
3 UPDATE ACCOUNT SET Bal = Bal + 100 WHERE AID = 2;
4 COMMIT;
```

### INITIALLY IMMEDIATE Example

```
1 ... DEFFERABLE INITIALLY IMMEDIATE ... /*Trigger*/
2 BEGIN TRANSACTION;
3 SET CONSTRAINTS bal_check_trigger DEFERRED
4 UPDATE ACCOUNT SET Bal = Bal - 100 WHERE AID = 1;
5 UPDATE ACCOUNT SET Bal = Bal + 100 WHERE AID = 2;
6 COMMIT;
```

- Multiple triggers defined for same event on the same table
- Order of trigger activation
  - BEFORE statement-level triggers
  - BEFORE row-level triggers
  - AFTER row-level triggers
  - AFTER statement-level triggers
- Within each category, triggers activated in alphabetic order
- If a BEFORE row-level trigger returns NULL, then subsequent triggers on same row are **omitted**.

### Functional Dependencies

A *normal form* is a definition of minimum requirements to reduce data redundancy and improve data integrity.

- An example of a FD can be: denoted as NRIC → Name.
- NRIC decides Name, if 2 records have the same NRIC, then they will always have the same name.
- To identify this: find a contradicting example: e.g. No 2 shops sell the same product:
- Contradicting example: Shop1 and Shop2 both sells ProductA, not allowed ∴  $P \rightarrow S$

### Armstrong’s Axioms

1. **Reflexivity:**  $AB \rightarrow A$
2. **Augmentation:** If  $A \rightarrow B$ , then  $A \rightarrow B$  and  $A \rightarrow C$  then  $AC \rightarrow BC$
3. **Transitivity:** If  $A \rightarrow B$ , and  $B \rightarrow C$  then  $A \rightarrow C$
4. **Decomposition:** If  $A \rightarrow BC$
5. **Union:** If  $A \rightarrow B$  and  $A \rightarrow C$

### Closure

Four attributes  $A, B, C, D$ . Given that  $B \rightarrow D, DB \rightarrow A, AD \rightarrow C$ , check if  $B \rightarrow C$  holds.

- $\{B\}^+ = \{B\}$
- $\{B\}^+ = \{B, D\}$ , since  $B \rightarrow D$
- $\{B\}^+ = \{B, D, A\}$ , since  $DB \rightarrow A$
- $\{B\}^+ = \{B, D, A, C\}$ , since  $AD \rightarrow C$
- $\{B, D, A, C\}$  referred to as *closure* of B, which is the set of componenets that can be ‘activated’ by B

Similarly to prove that something, e.g.  $X \rightarrow Y$  does not hold, we just need to show  $\{X\}^+$  **does not contain** Y.

### Superkeys / Keys of a table

**Superkeys:** set of attributes that decides all other attributes  
**Keys:** Superkey that is minimal. Whether or not a table has *redundancy* or *anomalies* depend on keys.

### Finding Keys

- Given a list of FDs, any attribute that *does not appear* on the RHS, it means that attribute **must be in every key!**
- Given a key: e.g.  $\{A, C\}$  is able to activate all other attributes, we do no need to check its supersets for *keys* since keys have to be minimal! e.g. no need to check  $\{A, C, B\}, \{A, C, D\}, \{A, C, B, D\}$  etc.

### Normal Forms

BCNF requires that if there is any **non-trivial** and **decomposed** FD  $A_1A_2...A_n \rightarrow B$ , then  $A_1A_2...A_n$  must be a superkey. More simply, BCNF requires LHS to be super keys, prevent redundancy.  
**Decomposed FD:** FD whose RHS only has 1 attribute  
**Non-Trival & Decomposed FD:** Decomposed FD whose RHS does not appear in LHS

1. Compute closure of each attribute subset
2. Derive keys of R (using closures)
3. Derive non-trivial and decomposed FDs from each closure
4. For each of the FD, if all LHS is super key then R satisfy BCNF

### Simplified BCNF

We have violation of BCNF iff we have a closure that satisfies the **more but not all** condition.

- Compute closure of each attribute subset
- If there is a closure (as below), then R is NOT in BCNF:
  - Closure contains some attributes not in  $\{A_1A_2...A_n\}$
  - Closure does not contain all attributes in table (more but not all closure)

### Normalization Algorithm

1. Find subset X of attributes in R, suhc that its closure  $\{X\}^+$  (i) does not contain more attributes than X, but (ii) does not contain all attributes in R
2. Decompose R to form 2 tables  $R_1$  and  $R_2$  such that
  - $R_1$  contains all attributes in  $\{X\}^+$
  - $R_2$  contains all attributes in X and attributes not in  $\{X\}^+$
3. If  $R_1$  nto in BCNF; further decompose  $R_1$ , same for  $R_2$

*BCNF decomposition not unique. If table only has 2 attributes, it MUST BE in BCNF: Check  $\emptyset$  set,  $A \rightarrow B, B \rightarrow A$  or both.*

### Projection of Closures/FDs

If we derive closures on table  $R_i$  decomposed from table R, we can enumerate attribute subsets of R. For each subset, derive its closure on R, thereafter project each closure onto R by removing those attributes that do not appear in  $R_i$ .

### Properties of BCNFs

No UPDATE/DELETE/INSERT anomalies. Small redundancy. Original table can be reconstructed from decomposed table.

- BCNF gurantees lossless join whenever common attributes in  $R_1$  and  $R_2$  contain a superkey of  $R_1$  or  $R_2$
- Decomposing R,  $R_1$  contains all attributes in  $\{X\}^+$ ,  $R_2$  contains all attributes in X as well as those not in  $\{X\}^+$
- Thus X will be the common attributes between both  $R_1$  and  $R_2$ , X will be a superkey of  $R_1$

However, dependencies may not be preserved.

## 3NF

A table satisfies 3NF if and only if for every non-trivial and decomposed FD:

- Either the LHS is **superkey**
- OR RHS is a **prime attribute** (appears in a key)
- 3NF more lenient than BCNF:
  - Satisfy BCNF  $\rightarrow$  satisfy 3NF (not necessarily vice versa)
  - Violate 3NF  $\rightarrow$  violate BCNF (not necessarily vice versa)

**FD Equivalence:** Prove that Given  $S$ ,  $S'$  can be derived, and vice versa.

## 3NF Check

1. Derive keys of table R
2. For each FD, check against LHS superkey OR each attribute on RHS is prime
3. If all FDs satisfy, then R is in 3NF

## 3NF Decomposition Algorithm

1. Given table R and set S of FDs, Derive minimal basis of S
2. In minimal basis, combine FDs whose LHS are the same
3. Create a table for each FD remained
4. If non of the tables contains a key of original table R, create table that contains a key of R (any key of R is ok)
5. Remove redundant tables

## Minimal Basis

Given set S of FDs, minimal basis of S, M is a simplified version of S, such that M equivalent to original set S.

1. every FD in S can be derived from M, and vice versa
2. every FD in M is a non-trivial and decomposed FD
3. if any FD is removed from M, then some FD in S cannot be derived from M
4. for any FD in M, if we remove an attribute from its left hand side, then the FD cannot be derived from S

## Adding Key for Lossless Join

$R(A, B, C, D)$ , with  $A \rightarrow B$ ,  $C \rightarrow D$

- Minimal basis:  $A \rightarrow B$ ,  $C \rightarrow D$
- Corresponding tables:  $R_1(A, B)$ ,  $R_2(C, D)$
- Notice that  $R_1$  and  $R_2$  cannot be used to reconstruct R
- This is why we require the following:
  - Check if any of the tables contain a key for R; if not, then create a table that contains a key for R
- In this case, R has only one key: AC
- Therefore, we add a table  $R_3(A, C)$

★ Use the key of original table to "guide" the JOIN back